Hortonworks

Apache Atlas: Data Governance

Partner Solutions

July 2015

Agenda

Overview

- Enterprise Goals
- Data Governance Initative

Atlas

- Feature tour
- Roadmap
- UI Tour

Demo

- Example: Sqoop
- Walk through step
- Search Tables / Tags

Hortonworks

Enterprise Data Governance Goals

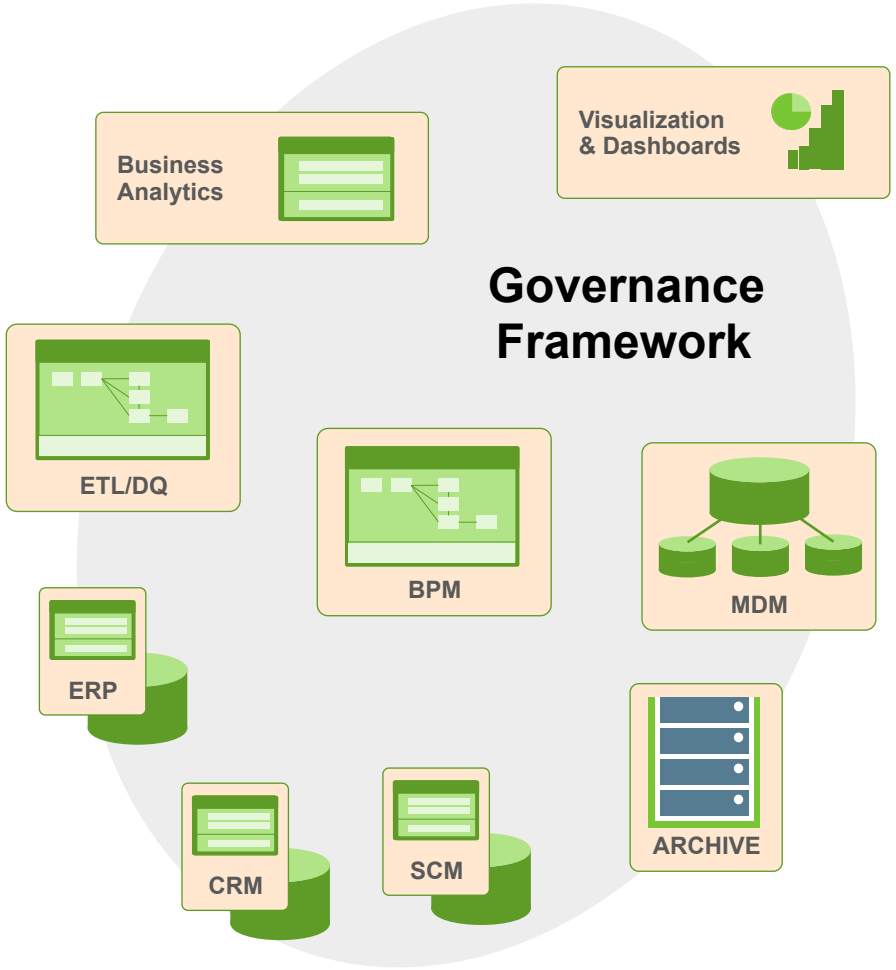

GOAL: Provide a common approach to data governance across all systems and data within the organization

- **Transparent**
  Governance standards & protocols must be clearly defined and available to all
- **Reproducible**
  Recreate the relevant data landscape at a point in time
- **Auditable**
  All relevant events and assets but be traceable with appropriate historical lineage
- **Consistent**
  Compliance practices must be consistent

Hortonworks

Data Governance Initiative for Hadoop

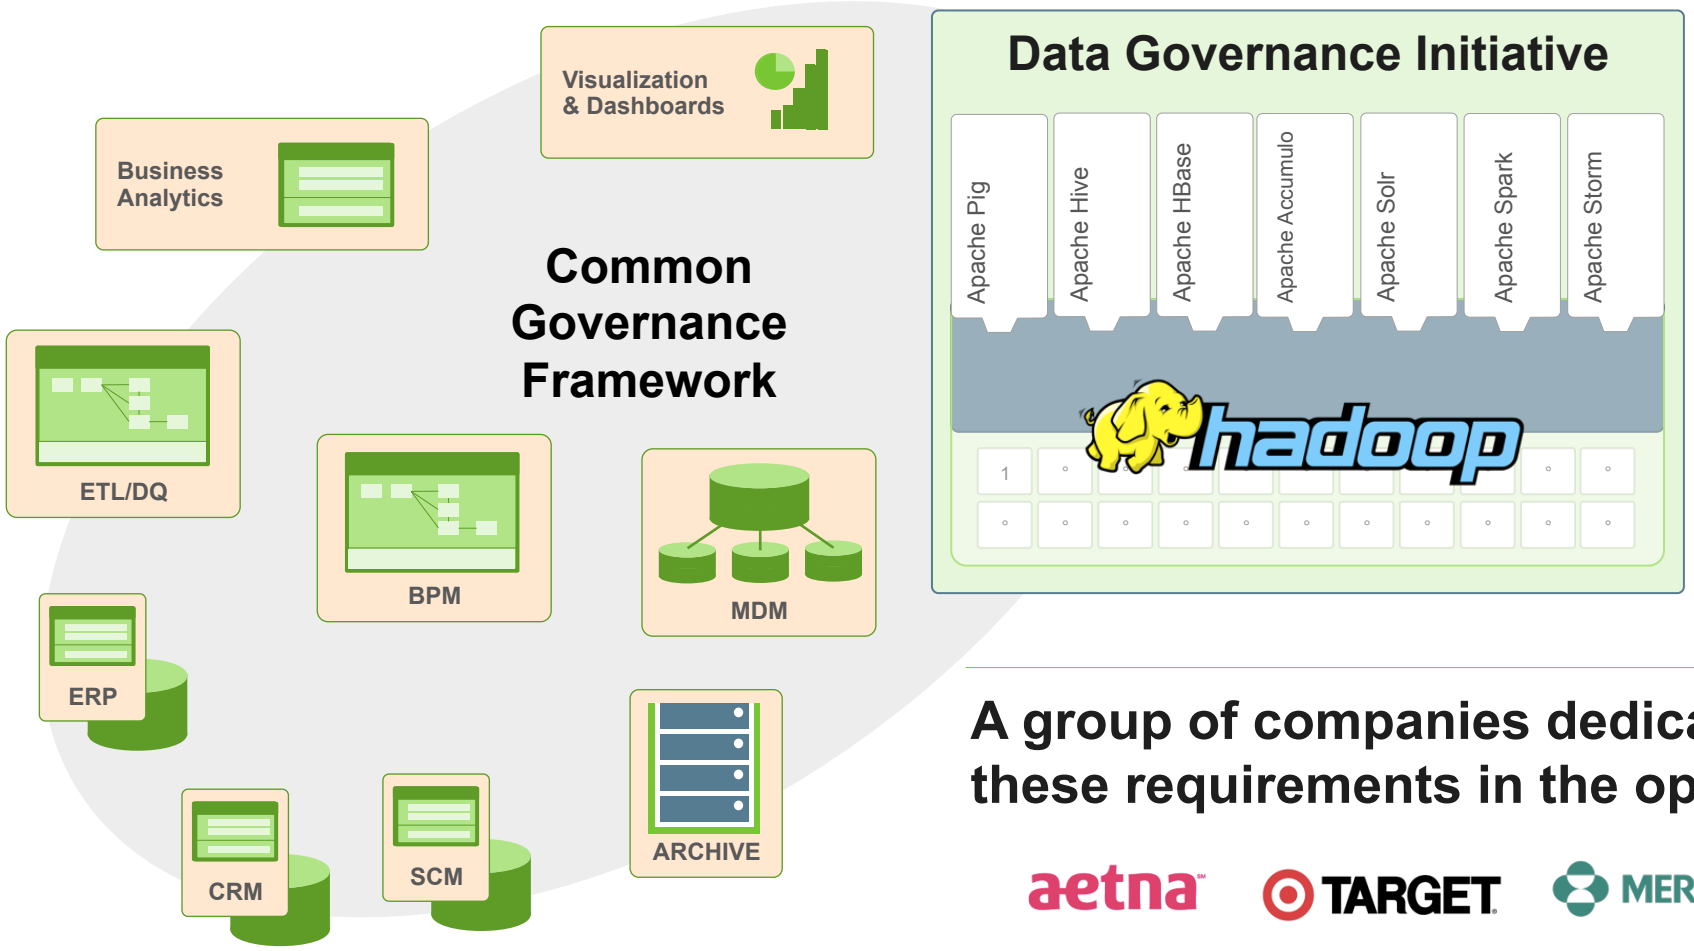

TWO Requirements

1. Hadoop must snap in to the existing frameworks and be a good citizen
2. Hadoop must also provide governance within its own stack of technologies

A group of companies dedicated to meeting these requirements in the open

aetna · TARGET · MERCK Be well · Hortonworks · SAP · Major Bank · Schlumberger · sas

Hortonworks

Apache Atlas Overview

We Do Hadoop

Apache Atlas Vision

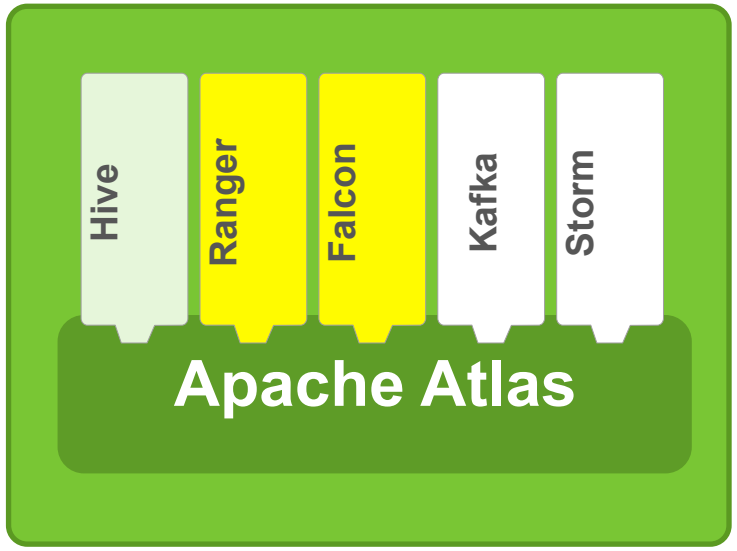

Metadata Services

- Flexible Knowledge Store
- Business Catalog / Operational Data
- Search & **Proscriptive Lineage**
- Centralized location for all metadata within HDP
- Interface point for Metadata Exchange with platforms outside of HDP.

Metadata will enrich every component

- Hive – Complete lineage, every HiveQL tracked
- Ranger – Tag or Attribute security ABAC
- Falcon – Business Taxonomy

Hortonworks

Apache Atlas Capabilities: Overview

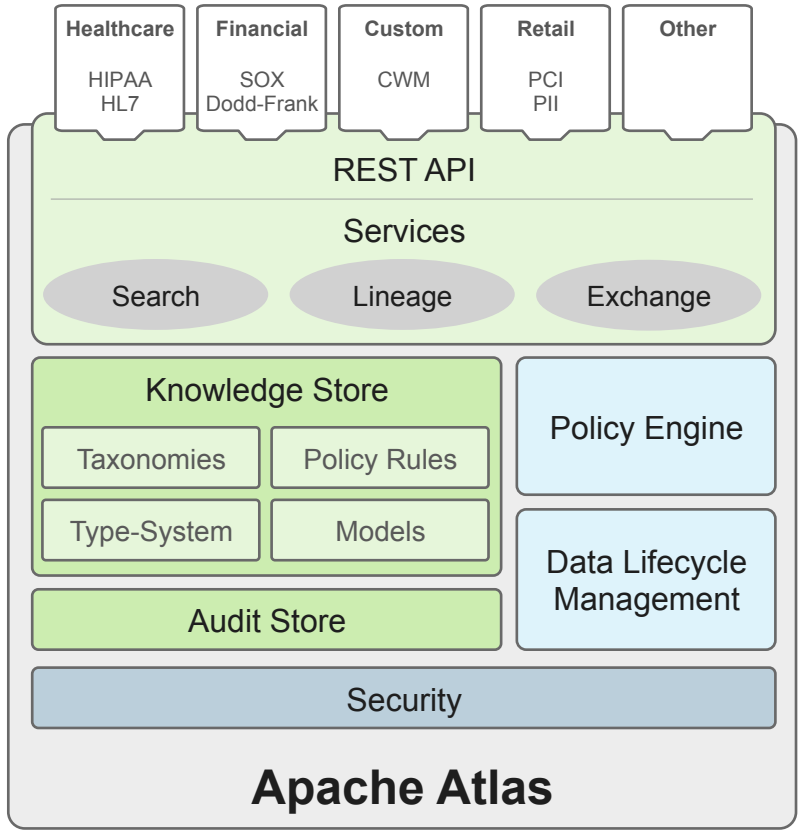

Data Classification

- Import or define taxonomy business-oriented annotations for data
- Define, annotate, and automate capture of relationships between data sets and underlying elements including source, target, and derivation processes
- Export metadata to third-party systems

Centralized Auditing

- Capture security access information for every application, process, and interaction with data
- Capture the operational information for execution, steps, and activities

Search & Lineage (Browse)

- Pre-defined navigation paths to explore the data classification and audit information
- Text-based search features locates relevant data and audit event across Data Lake quickly and accurately
- Browse visualization of data set lineage allowing users to drill-down into operational, security, and provenance related information

Security & Policy Engine

- Rationalize compliance policy at runtime based on data classification schemes
- Advanced definition of policies for preventing data derivation based on classification (i.e. re-identification)

Hortonworks

Sample Use Case: ETL Offload

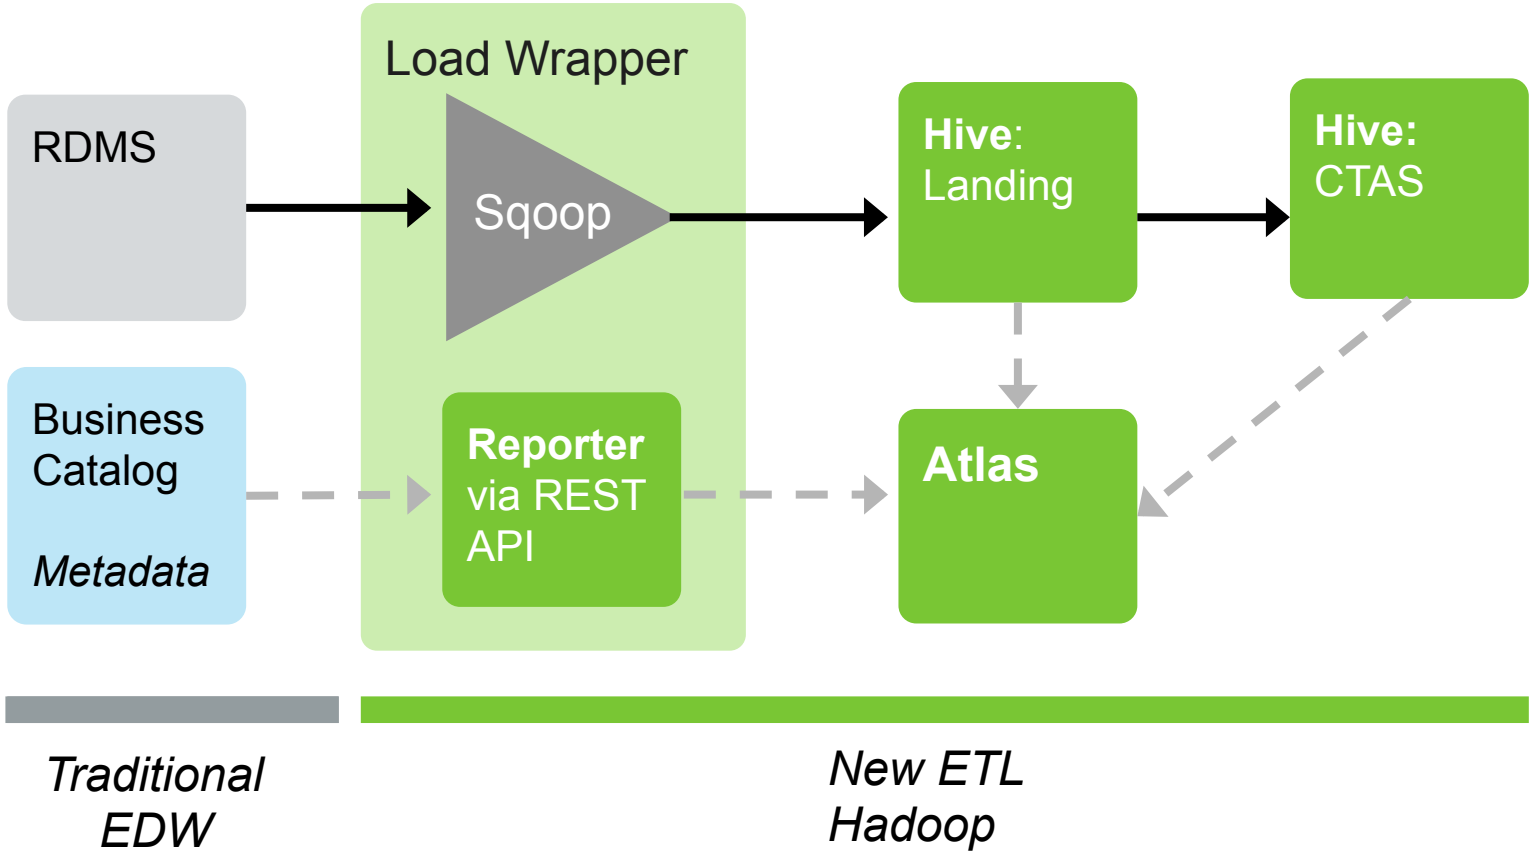

Hortonworks

Hive Integration

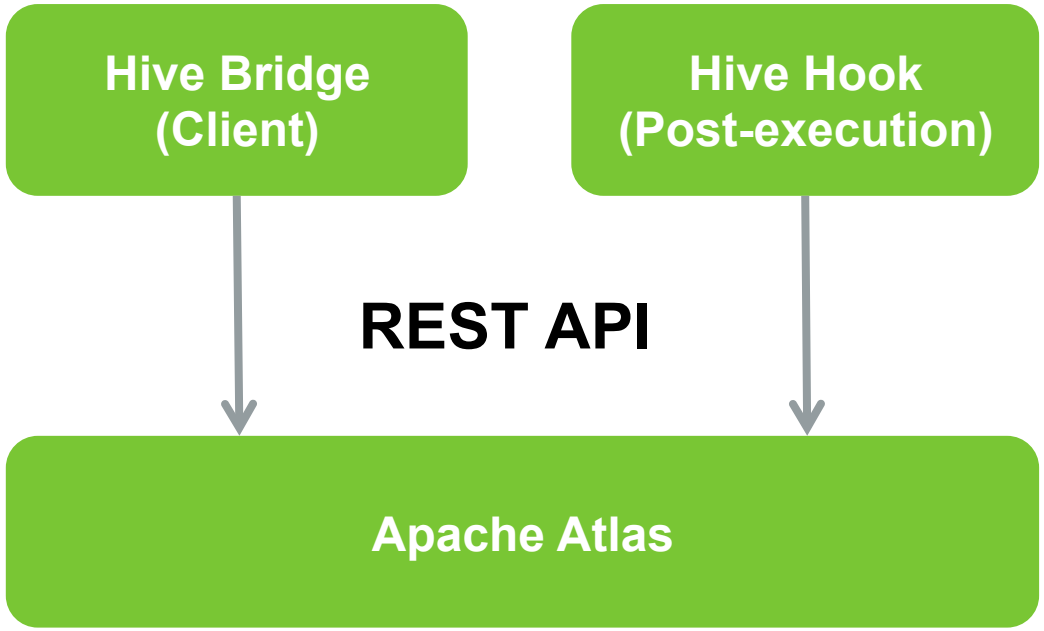

Governance Ready Certification Program

Curated group of vendor partners to provide rich & complete features

Customers choose features that they want to deploy – a la carte.

Low switching costs !

HDP at core to provide stability and interoperability

Hortonworks

High Level Roadmap

- **ASF MVP (May)** – Preview Core Metadata Services: Type system, API's, basic UI, Hive connecter
- **HDP 2.3 (July)** - GA Core Metadata Services. Preview Metadata Business Glossary
- **M10 – (Sept)** – Preview ABAC with Ranger integration and Preview Sqoop component connector
- **M20** – Preview Kafka, Storm connectors, Gov Ready Certification program, Preview row level & Column masking.
- **HDP 2.4 (Q4'15)** GA all preview features

Architecture

Hortonworks

High Level Architecture

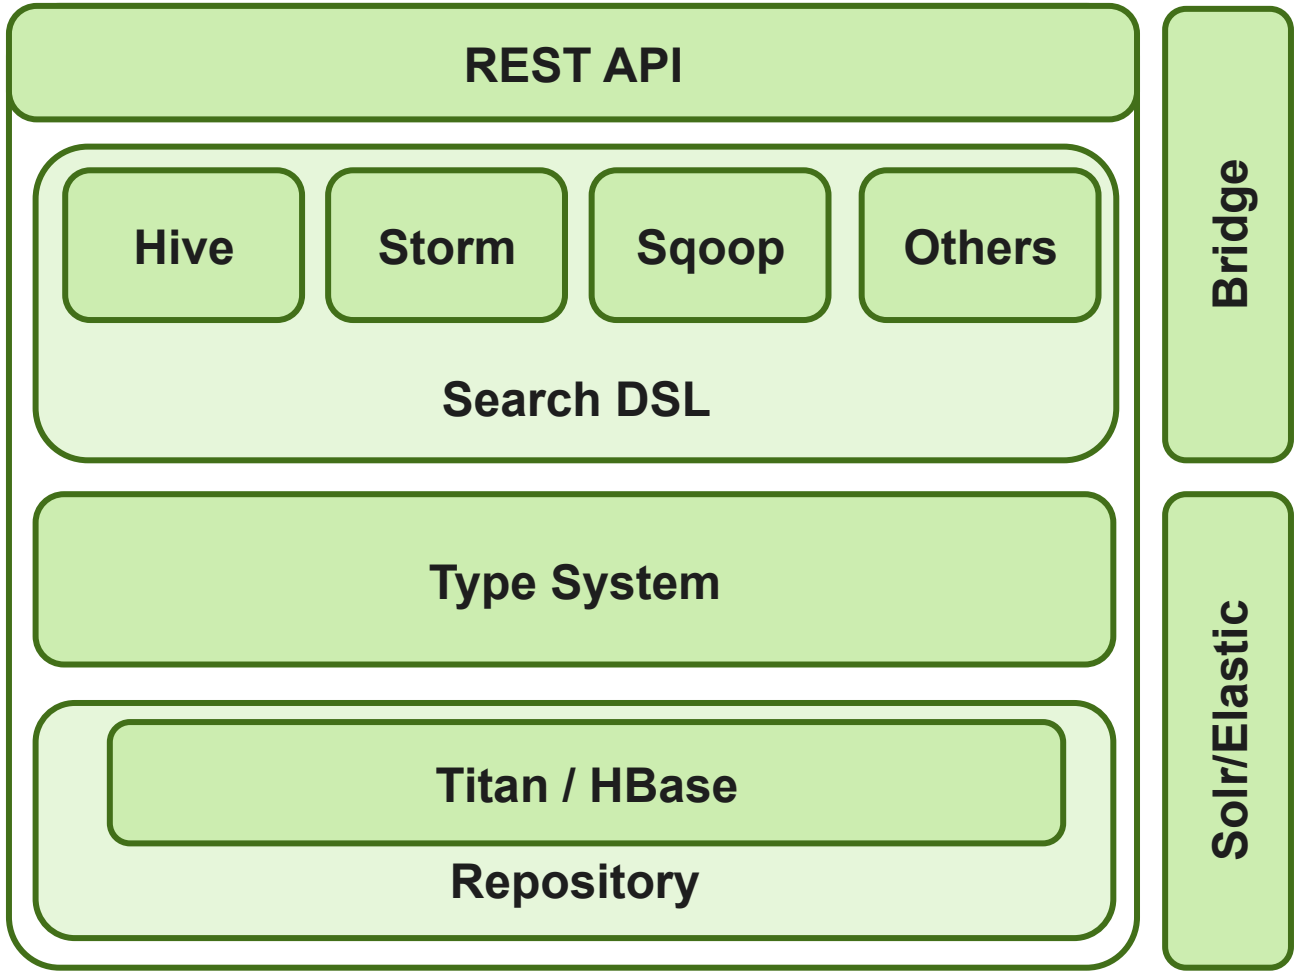

Hortonworks

Technology Stack

- Knowledge Store
  - Titan Graph DB
- Pluggable Search Backend
  - Elastic search
  - Solr
- Rules Engine
  - TBD

- Audit Store
  - YARN ATS - Time series DB
- Java 1.7
- Dashboard
  - TBD

APIs: Examples

Admin

GET: /admin/stack

GET: /admin/version

Entity

GET: /entities/definition/{guid}

POST: /entities/submit/{typeName}

GET: /entities/list/{entityType}

Metadata Discovery

GET: /discovery/search/gremlin/{gremlinQuery}

GET: /discovery/search/relationships/{guid}

GET: /discovery/search/fullText?text=<query>

GET: /discovery/getIndexedFields

Rexster

GET: /graph/vertices/{id}

GET: /graph/vertices/properties/{id}

GET: /graph/vertices

GET: /graph/vertices/{id}/{direction}

GET: /graph/edges/{id}

Types

POST: /types/submit/{typeName}

GET: /types/definition/{typeName}

GET: /types/list

Hive Lineage

GET: /bridge/hive/{id}

GET: /bridge/hive

POST: /bridge/hive

Hortonworks

Type System – Overview of Types

- Class
- Struct
- Trait
- Primitives

- Collections
  - Map
  - Array

- Instances (Entity)
  - Referenceable

Hortonworks

Type System – Data Types

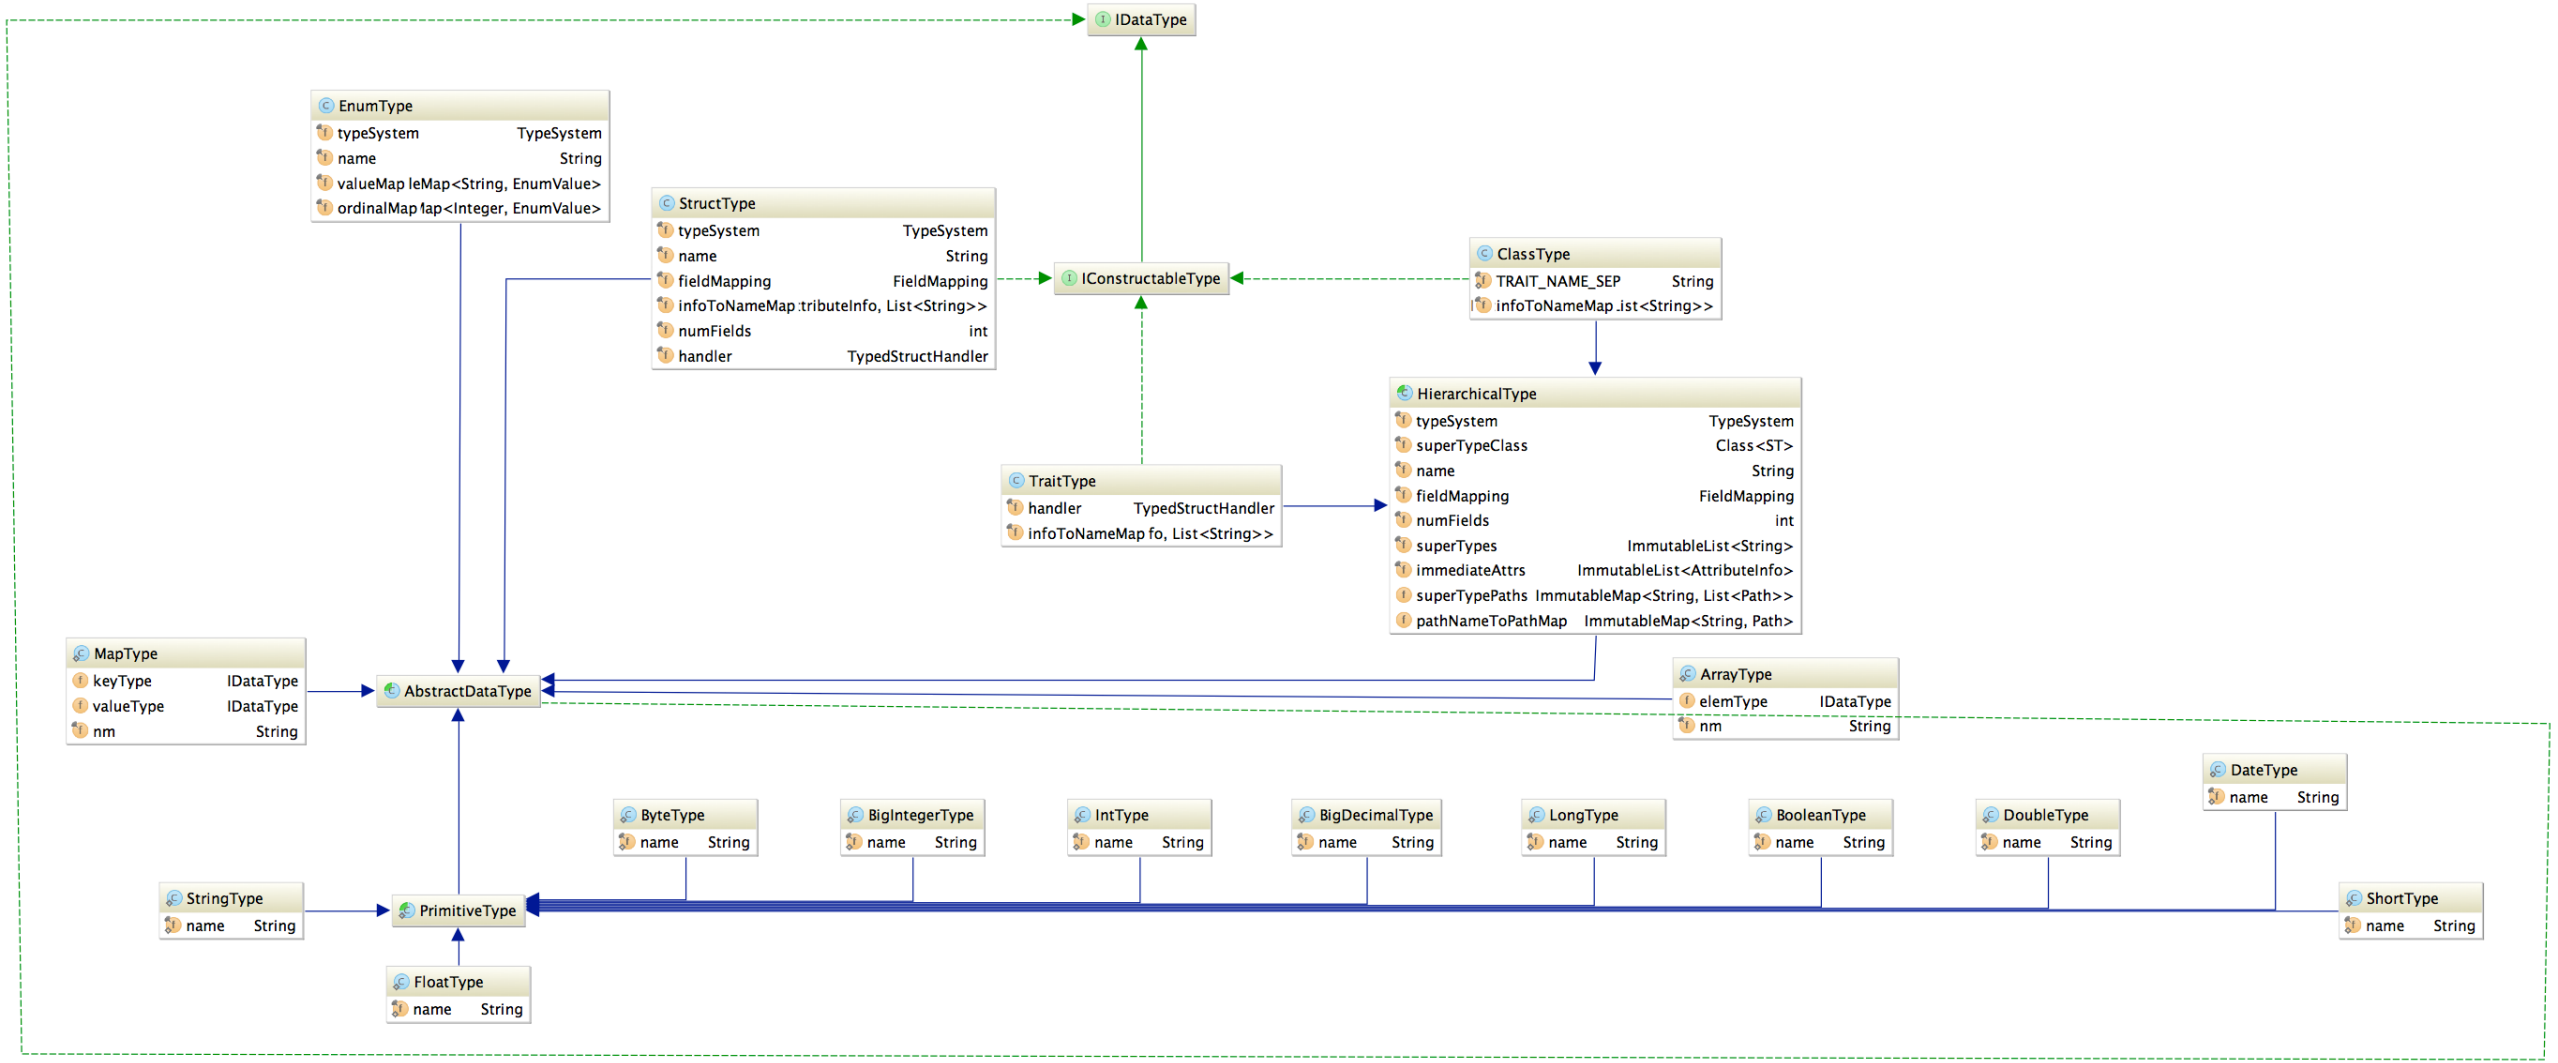

Hortonworks

```
_trait("Dimension") {}
_trait("PII") {}
_trait("Metric") {}
_trait("ETL") {}
_trait("JdbcAccess") {}

_class("DB") {
  "name" ~ (string, required,
indexed, unique)
  "owner" ~ (string)
  "createTime" ~ (int)
}

_class("StorageDesc") {
  "inputFormat" ~ (string,
required)
  "outputFormat" ~ (string,
required)
}
```

```
_class("Column") {
  "name" ~ (string, required)
  "dataType" ~ (string, required)
  "sd" ~ ("StorageDesc", required)
}


_class("Table", List()) {
  "name" ~ (string,  required,  indexed)
  "db" ~ ("DB", required)
  "sd" ~ ("StorageDesc", required)
}
```


Hortonworks

Repository

- **Graph Database**
  - Titan with storage backed by HBase
- **Types and instances are mapped to the Graph DB**
  - Classes, Structs and Traits map to a vertex
  - Relationships are mapped as edges
- **Search - plugin enabled**
  - Indexing based on type annotations
  - Solr
  - Elastic search

Hortonworks

Search

- **DSL with SQL Like Syntax**
  - from $type is $trait where $clause select|has $attributes loop $loopExpression withPath, repeat
- **Examples**
  - from DB
  - DB where name="Reporting" select name, owner
  - DB has name
  - DB is JdbcAccess
  - Column where Column is a PII
  - Table where name="sales_fact", columns
  - Table where name="sales_fact", columns as column select column.name, column.dataType, column.comment
- **Full-text search**

Hortonworks

Lineage

- **Uses Search DSL Loop expression**
  - Everything results in search
- **Named Queries**
- **inputs**
  - Table where (name = \"sales_fact_monthly_mv\") as src loop (LoadProcess->outputTable inputTables) as dest select src.name as src_name, dest.name as dest_name withPath
- **outputs**
  - Table where (name = \"sales_fact\") as src loop (LoadProcess->inputTables outputTables) as dest select src.name as src_name, dest.name as dest_name withPath
- **schema**
  - Table where name="sales_fact", columns

Hortonworks

Hive Integration

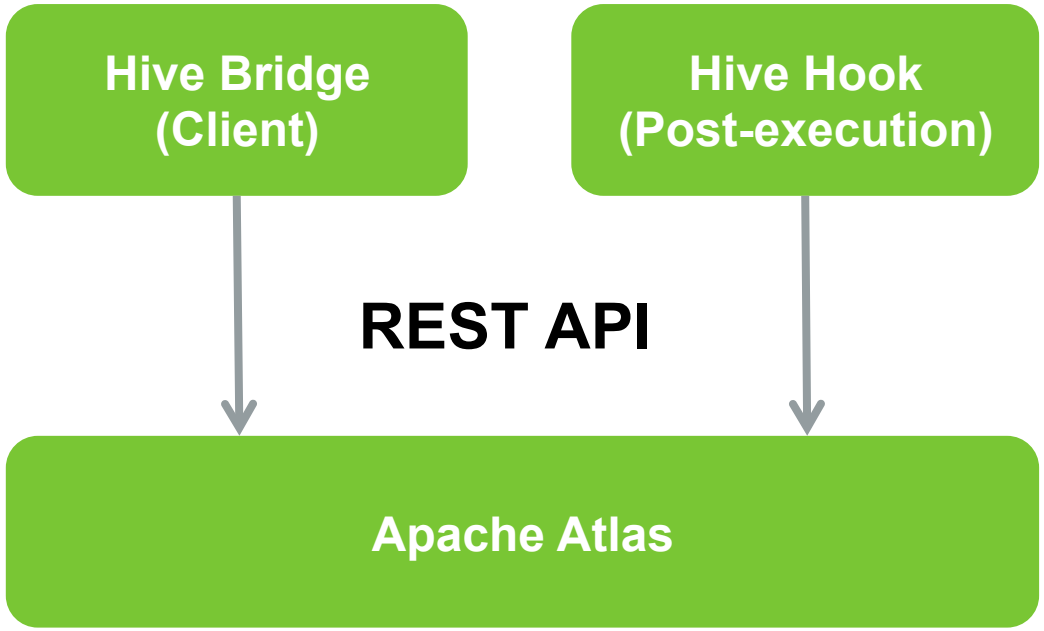

Apache Atlas Screens

Hortonworks

Apache Atlas

Table where name="sales_fact", columns

Search: Table, DB, Column

Tags

- Dimension
- ETL
- Fact
- JdbcAccess
- Metric
- PII

4 results matching your search query Table where name="sales_fact", columns were found

product_id

comment: product id , **dataType:** int

Tags :

time_id

comment: time id , **dataType:** int

Tags :

customer_id

comment: customer id , **dataType:** int

Tags : PII

sales

comment: product id , **dataType:** double

Tags : Metric

Previous 1 Next

Apache Atlas

PII

Search: Table, DB, Column

Tags

- Dimension
- ETL
- Fact
- JdbcAccess
- Metric
- PII

6 results matching your search query PII were found

619cc268-583e-4434-b2ce-0fd83a7f0e17
typeName: Column

2fa57bcd-ab08-4bf4-8743-17f36f4101ce
typeName: Column

9e70ca5d-59ba-4ed0-a5e1-c2d9de017605
typeName: Column

a58f3cba-c35e-4677-a8f9-2c5b43153220
typeName: Column

cbf03db4-d7f5-4c92-82cc-240330682222
typeName: Column

77c1af9c-75dd-4792-87ec-491732683654
typeName: Column

Apache Atlas

Back To Result

Name: sales_fact

Description: sales fact table

Details | Schema | Output | Input

Name	Comment	DataType
time_id	time id	int
product_id	product id	int
customer_id	customer id	int
sales	product id	double

Apache Atlas

Back To Result

Name: sales_fact

Description: sales fact table

Details | Schema | Output | Input

Key	Value
createTime	
db	id : 7791eb74-2b9a-4223-a676-0c612c17e68a jsonClass : org.apache.hadoop.metadata.typesystem.json.InstanceSerialization$_Id typeName : DB
lastAccessTime	
owner	Joe
retention	
sd	jsonClass : org.apache.hadoop.metadata.typesystem.json.InstanceSerialization$_Reference typeName : StorageDesc
tableType	Managed

Apache Atlas

Back To Result

Name: sales_fact_daily_mv

Description: sales fact daily materialized view

Details | Schema | Output | Input

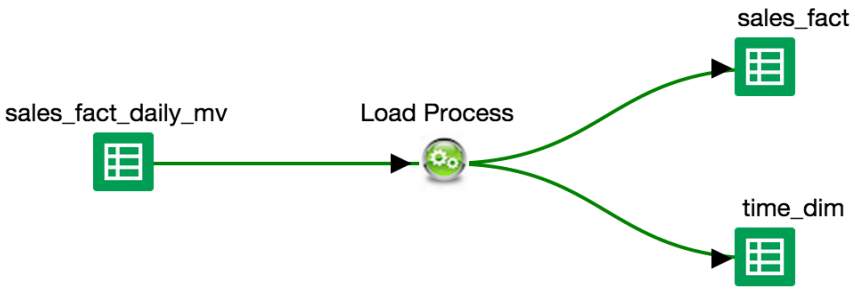

Demo Atlas

Atlas UI demostration

Search DSL

- Type – DB, Table, Column
- Tag - PII
- Keyword

Results

- Details
- Schema
- Lineage

Coming Features

Ingestion Demo Objective

- **Show Lineage with Sqoop Ingestion of data**
  - Custom process instrumention
- **Use the Hive Hook CTAS Operation**
  - Atlas Follow Lineage

- **Metadata Model in Atlas**
  - The Open Framework
  - Create Custom Types
  - Create Custom Process

- **Sample Codes**

Hortonworks

Setup

- **Source System**
  - MySQL Database
    - DRIVERS
    - TIMESHEET
  - Destination System
    - Single Node HDP 2.3 (Tech Preview)
      - Apache Atlas

Steps to Create Metadata

- **Create a Atlas Client Instance**
- **Create Type Definitions**
  - Class Types
  - Attributes
  - List the Types
- **Instantiate Entities**
  - - Create Entities (Class Type)
  - - Search the Types
- **Create Process**
  - Create DataSet Type
  - Create Process Type
  - Connect a Process Lineage

Attribute Definition

- **Name**
- **Data Type**
- **Multiplicity**
- **Composite**
- **isIndexable**
- **ReverseAttribute**

Questions and Answers

Hortonworks

Atlas Resources

- **HDP 2.3 Preview Sandbox VM:**
  - http://hortonworks.com/hdp/whats-new/
- **Apache Atlas:**
  - http://atlas.incubator.apache.org/
  - http://incubator.apache.org/projects/atlas.html
  - https://git-wip-us.apache.org/repos/asf/incubator-atlas.git
- **Partner Workshops**
  - http://hortonworks.com/partners/learn/
- **More to come with official GA release of HDP 2.3**

Thank you !

Hortonworks